

1. Introduction

Executing massive computational graphs on hardware with strict on-chip fast memory presents a combinatorial optimization challenge. Latency is governed by a roofline model that heavily penalizes repeated data spillage to slow memory. Our submission, the **V21 Engine**, minimizes end-to-end execution cycles through **Adaptive Dynamic Programming (DP)**, fine-grained **Split-K Pipelining**, strict memory limits enforced via **Deduplicated Retention Set** heuristics, **OpenMP parallelism over (w, h, k) granularity combos**, a **Step-Level Adaptive Window** that prevents timeouts on single-core systems, and **Operator Recomputation** that trades cheap recomputation for memory access savings.

The engine handles both trivial (5 nodes) and large (100+ nodes) DAGs without violating per-benchmark timeouts or causing OOM crashes. The binary is compiled statically with GCC and `-fopenmp`, deployed inside a Docker container based on Ubuntu 22.04. V21 achieves **39.85M total cycles** across all 5 public benchmarks (B1: 210K, B5: 545K, B9: 22.0M, B13: 11.1M, B17: 5.99M), with end-to-end wall-clock time well below the per-benchmark timeouts on 8 threads.




Figure 1: V21 Sliding Window DP Engine. The DAG is topologically sorted and processed via a sliding window; candidate partitions are evaluated in parallel across retention modes and granularity combos.

2. Dynamic Programming Core

The core challenge is grouping dependent operations into subgraphs without saturating the fast memory capacity C_{fast} . Greedy algorithms fall into local optima — aggressively retaining intermediate tensors often leads to OOM crashes deeper in the graph. RL approaches require exorbitant training and cannot guarantee generalization within 10 minutes.

We formulate the problem as linear **Dynamic Programming** over a topologically sorted node sequence. The state $dp[i]$ represents the minimal total latency to execute the graph up to node i :

$$dp[i] = \min_{j < i} \{ dp[j] + L(S_{j \rightarrow i-1}, M_{\text{retention}}) \}$$

where L calculates the analytical roofline latency, parameterized by the tensors retained in fast memory from the preceding step.

2.1 Adaptive Window Sizing

Exhaustively evaluating all $j \in [0, i-1]$ is $O(N^2)$, intractable for dense $N \geq 60$ topologies. V21 uses **Adaptive Windowing** with two window levels per size class: a wide **W_max** when the system runs fast (8-core hardware), and a conservative **W_cons** as fallback on 1-core systems to prevent timeouts:

N	W_max	W_cons	nm_max	nm_cons
≤ 32	N (full)	N (full)	10	10
≤ 40	24	20	8	6
≤ 65	20	12	8	5
≤ 105	16	10	6	5
> 105	9	7	4	4

The dual-window design (Section 2.2) lets V21 use aggressive widths on fast hardware while gracefully degrading to V19-equivalent work on slow hardware — no timeouts, no regressions.

2.2 Step-Level Adaptive Fallback

A per-step timer tracks elapsed time after each outer DP step i . If the current step’s wall-clock cost multiplied by the remaining steps would exceed the available time budget, the engine **permanently** switches to $(W_{\text{cons}}, nm_{\text{cons}})$ for all subsequent steps. This step-level adaptation is more robust than progress-extrapolation: early DP steps are cheap (small j -range) and do not predict the cost of later steps at full window width. The fallback triggers before any timeout risk materializes; on 8-core hardware it never triggers at all.

3. Memory Hierarchy & Split-K Tiling




Figure 2: Memory Hierarchy & Split-K Tiling. The scheduler decides SRAM vs. DRAM placement; Split-K partitions the K dimension keeping partial sums pinned in fast memory.

3.1 Granularity Candidate Generation

Once a candidate subgraph S is identified, determining its execution tile dimensions $[w, h, k]$ is the physical optimization lever. V21 curates candidates from: (1) multiples/divisors of the native hardware granularity, (2) dimensional bounds of topological outputs, (3) power-of-2 factors, (4) divisors of tensor dimensions, and (5) **memory-aware candidates** computed from available fast memory slack. Memory-aware combos are now generated for problems up to

$N < 65$ (extended from $N < 32$ in V15). All candidates are flattened into a single vector and evaluated in parallel (Section 7).

3.2 Split-K Pipelining for MatMul

MatMul operations require holding LHS, RHS, and Output simultaneously. V21 evaluates the inner reduction depth k via **Split-K Pipelining**: using small k values, the system accumulates partial dot products within a pinned Output tensor in fast memory, iterating over K sequentially. This slashes the working set requirement, bypassing OOM crashes in large GEMMs without forfeiting computational efficiency.

4. Fast Memory Retention Strategies

Deciding which tensors to persist in fast memory is NP-Hard — tensors vary in size and future utility. Each DP transition generates $M_{\text{retention}}$ by simulating up to **10 concurrent heuristics**:

- **Mode 0 (Size, large first)**: Greedy pack by descending size, full budget.
- **Mode 1 (Value-based)**: Score = size / distance-to-next-use. Large tensors needed soon get priority. First actual forward use is scanned to avoid $\text{lu}[t]$ overestimation.
- **Mode 2 (Proximity, 3/4 budget)**: Nearest-use within 8 ops; 75% capacity cap.
- **Mode 3 (Chain-aware)**: Sort by precomputed `total_fanout` descending. Skip/residual tensors consumed by many future ops are retained first.
- **Mode 4 (lu-distance)**: Sort by `last_use` ascending — tensors that expire soonest are prioritized; they must be in fast memory for the next few subgraphs and won't be needed long after.
- **Modes 5–9**: Size-small-first threshold, wider proximity windows (lookahead 12), smallest-first full budget, near-only 8-op window, half-budget narrow lookahead.

Hash Deduplication. Generating 10 retention sets produces overlapping configurations. V21 computes a 64-bit FNV-1a hash per set — only unique hashes proceed to full evaluation, eliminating 30–50% redundant work.

5. V21 Operator Recomputation

V21 introduces **Operator Recomputation** as a novel extension of the DP transition. For small subgraphs ($|S| \leq 4$ ops), the engine identifies candidate input tensors that:

- Are consumed by the subgraph but not produced within it.
- Are **not** in the retained set from the prior step.
- Are **not** graph-level inputs (must have a known producer).
- Have `last_use[t] <= sg_last` — the tensor is fully consumed within this subgraph and not needed afterward; tensors needed in later subgraphs are excluded.
- Are **not** graph-level outputs.
- Have a Pointwise producer P before j whose recompute cost satisfies:

$$\text{base_cost}(P) \times \text{slow_bw} < \frac{\text{tsize}(t)}{2}$$

i.e., recomputing P is cheaper than half the cost of loading tensor t from slow memory. When such candidates exist, the engine evaluates **both** the base subgraph and the **extended** subgraph (with the cheap Pointwise ops prepended in topological order). The cheaper result wins — guaranteeing no regressions while capturing fusion opportunities that eliminate slow-memory loads.

A `producer_of[]` array (built once at solve-start, `producer_of[t] = op index producing t`) enables O(1) producer lookup during transition evaluation. Recomputed op indices are stored alongside the standard subgraph and emitted first in the output traversal order.

6. Latency Evaluation Model

The evaluation function simulates hardware analytically. For every candidate, latency is bounded by:

$$L = \max(T_{\text{compute}}, T_{\text{stream_in}} + T_{\text{stream_out}})$$

where T_{compute} is the arithmetic cost / compute throughput, and T_{stream} is bytes transferred / memory bandwidth. The model accounts for tile-level working sets, distinguishing streaming tensors (loaded per tile) from stationary tensors (loaded once and reused).

Traversal Order. V21 evaluates **Row-Snake** (reuses LHS row-by-row) and **Col-Snake** (reuses RHS column-by-column) patterns. The traversal with lower total stream bandwidth is selected. For single-tile subgraphs, no order is emitted.

7. Concurrency Architecture

OpenMP Granularity Parallelism. Inside `find_best()`, all (w, h, k) candidate combos are flattened into a single vector and distributed across threads via `#pragma omp parallel + #pragma omp for schedule(dynamic, 4)`, with each thread maintaining a local `Best` and merging under `#pragma omp critical`. V21 uses a single level of parallelism (`omp_set_max_active_levels(1)`); retention modes are iterated sequentially in the outer loop, relying on hash deduplication (Section 4) to skip redundant work rather than spawning additional threads. This keeps scheduling overhead low on the 8-core target while still exposing hundreds–thousands of independent combo evaluations per DP transition.

Cache-Optimized Structures. (1) **Flat Bitset (TensorSet)**: 8192-bit bitset (64-bit words); O(1) insert/lookup; fits in L1 cache. (2) **Precomputed SubgraphTensorInfo**: tensor metadata computed once per DP transition, shared across all (w, h, k) and retention-mode evaluations. (3) **OpType Enum**: runtime string comparisons replaced by an integer enum parsed at load time.

8. Results

Test environment. Windows 11 host with WSL2 running Ubuntu 22.04; Intel Core i7-8750H (6 cores / 12 threads @ 2.20 GHz), 16 GB RAM (≈ 8 GB exposed to WSL); compiled with `g++ 11.4.0 -O3 -std=c++20 -fopenmp`. Benchmarks were run inside WSL with `OMP_NUM_THREADS` set explicitly to 1 and 8.

Benchmark	Cycles	Time (T8)	Time (T1)	Timeout
B1 (5 ops)	210,219	0.16s	0.02s	2s
B5 (19 ops)	545,348	0.09s	0.33s	5s
B9 (32 ops)	22,028,095	2.48s	16.53s	15s
B13 (63 ops)	11,078,205	8.51s	26.52s	30s
B17 (103 ops)	5,992,359	4.74s	27.69s	60s
Total	39,854,226	15.97s	71.09s	—

Cycles are reported directly by V21's internal roofline model and are essentially independent of thread count (differences below 10^{-4} relative across T1 vs. T8). With 8 threads (the official evaluation setup), V21 stays well inside every per-benchmark timeout — total wall-time is 16s against a combined 112s budget. On a single thread, B9 exceeds its 15s budget by $\{\approx\}1.5$ s; the step-level adaptive fallback (Section 2.2) is what keeps B13 and B17 comfortably below their 30s and 60s limits respectively, by switching to the conservative window as soon as extrapolated cost threatens the budget.

The V21 engine merges DP optimality bounds with hyper-localized heuristics for tiling, memory pinning, and now operator recomputation. Step-level adaptive windowing guarantees timeout-safety on the 8-core deployment target and degrades gracefully on single-core hardware. Hash-deduplicated retention sets, bitset-optimized hot paths, and precomputed tensor metadata collectively minimize per-transition search overhead. Pointwise recomputation further reduces slow-memory traffic for small subgraphs with cheap producers.